

UNICORE - From Project Results to Production Grids

A. Streit, D. Erwin, Th. Lippert, D. Mallmann, R. Menday, M. Rambadt, M. Riedel,
M. Romberg, B. Schuller, and Ph. Wieder

John von Neumann-Institute for Computing (NIC)

Forschungszentrum Jülich (FZJ)

52425 Jülich, Germany

E-mail: {a.streit, d.erwin, th.lippert, d.mallmann, r.menday, m.riedel,
m.rambadt, m.romberg, b.schuller, ph.wieder}@fz-juelich.de

The UNICORE Grid-technology provides a seamless, secure and intuitive access to distributed Grid resources. In this paper we present the recent evolution from project results to production Grids. At the beginning UNICORE was developed as a prototype software in two projects funded by the German research ministry (BMBF). Over the following years, in various European-funded projects, UNICORE evolved to a full-grown and well-tested Grid middleware system, which today is used in daily production at many supercomputing centers worldwide. Beyond this production usage, the UNICORE technology serves as a solid basis in many European and International research projects, which use existing UNICORE components to implement advanced features, high level services, and support for applications from a growing range of domains. In order to foster these ongoing developments, UNICORE is available as open source under BSD licence at SourceForge, where new releases are published on a regular basis. This paper is a review of the UNICORE achievements so far and gives a glimpse on the UNICORE roadmap.

1. Introduction

End of 1998 the concept of “Grid computing” was introduced in the monograph “The Grid: Blueprint for a New Computing Infrastructure” by I. Foster and C. Kesselman [1]. Two years earlier, in 1997, the development of the UNICORE - Uniform Interface to Computing Resources - system was initiated to enable German supercomputer centers to provide their users with a *seamless, secure, and intuitive access* to their heterogeneous computing resources. Like in the case of the Globus Toolkit® [2] UNICORE was started before “Grid Computing” became the accepted new paradigm for distributed computing.

The UNICORE vision was proposed to the German Ministry for Education and Research (BMBF) and received funding. A first prototype was developed in the UNICORE¹ project [3]. The foundations for the current production version were laid in the follow-up project UNICORE Plus² [4], which was successfully completed in 2002. Since then UNICORE was used in operation at German supercomputing centers and became a solid

¹funded by BMBF grant 01 IR 703, duration: August 1997 - December 1999

²funded by BMBF grant 01 IR 001 A-D, duration: January 2000 - December 2002

basis for numerous European projects. In this paper we will describe the evolution of UNICORE from a prototype software developed in research projects to a Grid middleware used today in the daily operation of production Grids.

Although already set out in the initial UNICORE project proposal in 1997, the goals and objectives of the UNICORE technology are still valid:

- Foremost, the aim of UNICORE is to hide the rough edges resulting from different hardware architectures, vendor specific operating systems, incompatible batch systems, different application environments, historically grown computer center practices, naming conventions, file system structures, and security policies – just to name the most obvious.
- Equally, security is a constituent part of UNICORE’s design relying on X.509 certificates for the authentication of users, servers, and software, and the encryption of the communication over the internet.
- Finally, UNICORE is usable by scientists and engineers without having to study vendor or site-specific documentation. A Graphical User Interface (GUI) is available to assist the user in creating and managing jobs.

Additionally, several basic conditions are met by UNICORE: the Grid middleware supports operating systems and batch systems of all vendors present at the partner sites. In 1997 these were for instance large Cray T3E systems, NEC and Hitachi vector machines, IBM SP2s, and smaller Linux clusters. Nowadays the spectrum is even broader, of course with modern hardware, such as IBM p690 systems. The deployed software has to be non-intrusive, so that it does not require changes in the computing centers hardware and/or software infrastructure. Maintaining site autonomy is still a major issue in Grid computing, when aspects of acceptability and usability in particular from the system administrator’s point of view are addressed. In addition to UNICORE’s own security model, site-specific security requirements (e. g. firewalls) are supported.

Near the end of the initial funding period of the UNICORE Plus project, a working prototype was available, which showed that the initial concept works. By combining innovative ideas and proven components over the years, this first prototype evolved to a *vertically integrated* Grid middleware solution.

The remainder of this paper is structured as follows. In Section 2 the usage of UNICORE in production is presented to give potential users a high-level glimpse on the UNICORE technology and architecture. In section 3 UNICORE’s architecture and core features are described in more detail. European funded projects, which use UNICORE as a basis for their work are described in Section 4. Section 5 gives an outlook on the future development of UNICORE. The paper closes with conclusions and acknowledgements.

2. UNICORE in Production

From its birth in two German BMBF-funded projects to its extensive use and further development in a variety of EU and BMBF research projects (cf. Section 4 for details), the UNICORE technology ran through an evolutionary process transforming from an initial prototype software to a powerful production Grid middleware.

2.1. UNICORE@SourceForge

Since May 2004, the UNICORE technology with all its components is available as open source software under the BSD license. It can be downloaded from the SourceForge repository. Besides the core developers of UNICORE (namely Fujitsu Laboratories of Europe, Intel Germany and the Research Center Jülich), there are numerous contributors from all over the world, e.g. Norway, Poland, China and Russia. The Web site [5] offers a convenient entry point for interested users and developers. In the download section the UNICORE software is bundled in different packages, e.g. the client package and individual packages for the different server components Gateway, NJS, TSI/IDB, UUDB (cf. Section 3 for details), and plug-ins. Until January 2005 more than 2800 downloads of UNICORE are counted.

A tracker section linked on the Web site establishes a communication link to the core developer community. The corresponding mailing lists allow users to report bugs, to request new features, and to get informed about bug fixes or patches. For the announcement of new software releases a separate mailing list was created. The Grid team at the Research Center Jülich is responsible for UNICORE@SourceForge. Its work includes coordinating and driving the development effort, and producing consolidated, stable, and tested releases of the UNICORE software.

2.2. Production System on Jump

Since July 2004 UNICORE is established as production software to access the supercomputer resources of the John von Neumann-Institute for Computing (NIC) at the Research Center Jülich. These are the 1312-processor IBM p690 cluster (Jump) [6], the Cray SV1 vector machine, and a new Cray XD1 cluster system. As an alternative to the standard SSH login, UNICORE provides an intuitive and easy way for submitting batch jobs to the systems. The academic and industrial users come from all over Germany and from parts of Europe. The applications come from a broad field of domains, e.g. astrophysics, quantumphysics, medicine, biology, chemistry, and climate research, just to name the largest user communities. A dedicated, pre-configured UNICORE client with all required certificates and accessible Vsites is available for download. This alleviates the installation and configuration process significantly. Furthermore, an online installation guide including a certificate assistant, an user manual, and example jobs help users getting started.

To provide the NIC-users with adequate certificates and to ease the process of requesting and receiving a certificate, a certificate authority (CA) was established. User certificate requests are generated in the client and have to be send to the CA. Since introduction of UNICORE at NIC, more than 120 active users requested a UNICORE user certificate.

A mailing list serves as a direct link of the users to UNICORE developers in the Research Center Jülich. The list allows to post problems, bug reports, and feature requests. This input is helpful in enhancing UNICORE with new features and services, in solving problems, identifying and correcting bugs, and influences new releases of UNICORE available at SourceForge.

2.3. DEISA – Distributed European Infrastructure for Scientific Applications

Traditionally, the provision of high performance computing resources to researchers has traditionally been the objective and mission of national HPC centers. On the one hand, there is an increasing global competition between Europe, USA, and Japan with

growing demands for compute resources at the highest performance level, and on the other hand stagnant or even shrinking budgets. To stay competitive major investments are needed every two years – an innovation cycle that even the most prosperous countries have difficulties to fund.

To advance science in Europe, eight leading European HPC centers devised an innovative strategy to build a Distributed European Infrastructure for Scientific Applications (DEISA)³ [7]. The centers join in building and operating a tera-scale supercomputing facility. This becomes possible through deep integration of existing national high-end platforms, tightly coupled by a dedicated network and supported by innovative system and grid software. The resulting virtual distributed supercomputer has the capability for natural growth in all dimensions without singular procurements at the European level. Advances in network technology and the resulting increase in bandwidth and lower latency virtually shrink the distance between the nodes in the distributed super-cluster. Furthermore, DEISA can expand horizontally by adding new systems, new architectures, and new partners thus increasing the capabilities and attractiveness of the infrastructure in a non-disruptive way.

By using the UNICORE technology, the four core partners of the projects have coupled their systems using virtually dedicated 1 Gbit/s connections. The DEISA super-cluster currently consists of over 4000 IBM Power 4 processors and 416 SGI processors with an aggregated peak performance of about 22 teraflops. UNICORE provides the seamless, secure and intuitive access to the super-cluster.

The Research Center Jülich is one of the DEISA core partners and is responsible for introducing UNICORE as Grid middleware at all partner sites and for providing support to local UNICORE administrators.

In the following we describe the DEISA architecture. Note, a detailed description of UNICORE's architecture and server components can be found in Section 3 and in particular in Figure 2. All DEISA partners have installed the UNICORE server components Gateway, NJS, TSI, and UUDB to access the local supercomputer resources of each site via UNICORE. Figure 1 shows the DEISA UNICORE configuration. For clarity only four sites are shown. At each site, a Gateway exists as an access to the DEISA infrastructure. The NJSs are not only registered to their local Gateway, but to all other Gateways at the partner sites as well. Local security measures like firewall configurations need to consider this, by permitting access to all DEISA users and NJSs. This fully connected architecture has several advantages. If one Gateway has a high load, access to the high performance supercomputers through DEISA is not limited. Due to the fully connected architecture, no single point of failure exists and the flexibility is increased.

The DEISA partners operate different supercomputer architectures, which are all accessible through UNICORE. Initially all partners with IBM p690 clusters are connected to one large virtual supercomputer. In a second step other supercomputers of different variety are connected to DEISA, making the virtual supercomputer heterogeneous. UNICORE can handle this, as it is designed to serve such heterogeneous architectures in a seamless, secure, and intuitive way.

In December 2004 a first successful UNICORE demonstration between the four DEISA

³funded by EC grant FP6-508803, duration: May 2004 - April 2009

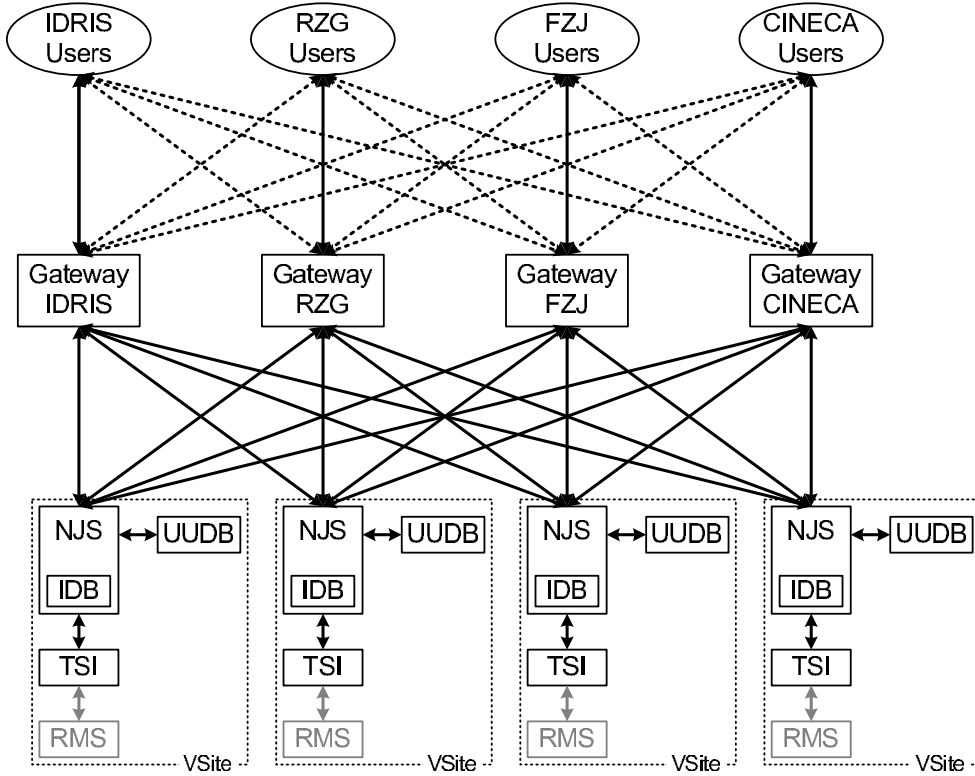


Figure 1. The DEISA architecture.

core sites FZJ (Research Center Jülich, Germany), RZG (Computing Center Garching, Germany), CINECA (Italian Interuniversity Consortium, Italy) and IDRIS (Institute for Development and Resources in Intensive Scientific Computing, France) was given. Different parts of a distributed astrophysical application were generated and submitted with UNICORE to all four sites.

The experience and knowledge of the researchers, developers, users, and administrators in working with UNICORE in the DEISA project on a large production platform will be used as useful input for future developments of the UNICORE technology. A close synchronization with the UniGrids project (cf. Section 5.1) is foreseen.

3. The Architecture of UNICORE

Figure 2 shows the layered Grid architecture of UNICORE consisting of user, server and target system tier [8]. The implementation of all components shown is realized in Java. UNICORE meets the Open Grid Services Architecture (OGSA) [9] concept following the paradigm of 'Everything being a Service'. Indeed, an analysis has shown that the basic ideas behind UNICORE already realizes this paradigm [10,11].

3.1. User Tier

The UNICORE Client provides a graphical user interface to exploit the entire set of services offered by the underlying servers. The client communicates with the server tier by sending and receiving Abstract Job Objects (AJO) and file data via the UNICORE Protocol Layer (UPL) which is placed on top of the SSL protocol. The AJO is the realization of UNICORE's job model and central to UNICORE's philosophy of abstraction and seamlessness. It contains platform and site independent descriptions of computational and data related tasks, resource information and workflow specifications along with user and security information. AJOs are sent to the UNICORE Gateway in form of serialized and signed Java objects, followed by an optional stream of bytes if file data is to be transferred.

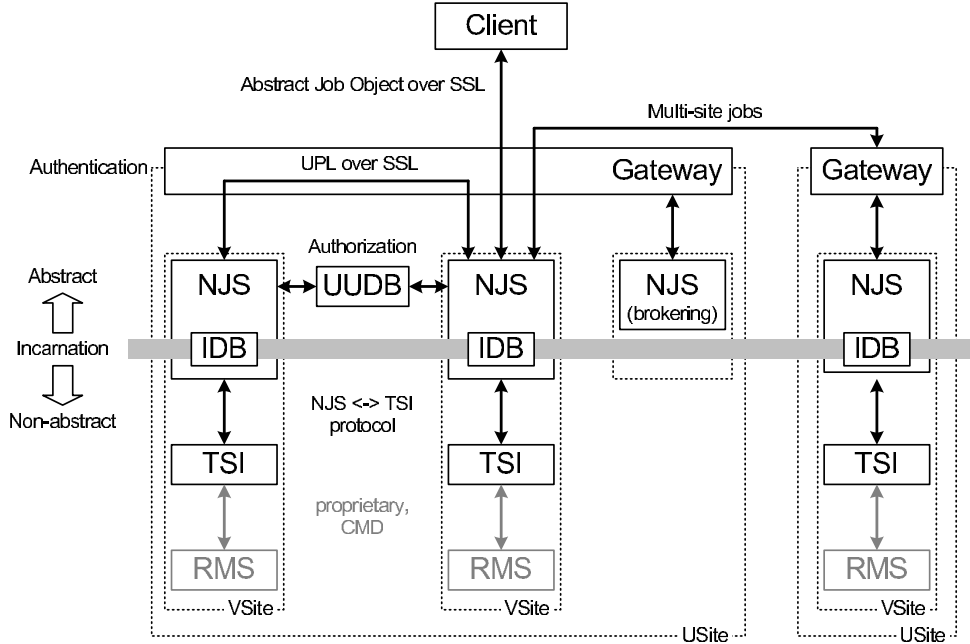


Figure 2. The UNICORE architecture.

The UNICORE client assists the user in creating complex, interdependent jobs that can be executed on any UNICORE site (U-site) without requiring any modifications. A UNICORE job, more precisely a job group, may recursively contain other job groups and/or tasks and may also contain dependencies between job groups to generate job workflows. Besides the description of a job as a set of one or more directed a-cyclic graphs, conditional and repetitive execution of job groups or tasks are also included. For the monitoring of jobs, their status is available at each level of recursion down to the individual task. Detailed log information is available to analyze potential error conditions. At the end of the execution of the job it is possible to retrieve the **stdout** and **stderr** output of the job. Data management functions like import, export, and transfer are available

through the GUI as explicit tasks. This allows the user to specify data transfer from one target system to another (e.g. for workflows), from or to the local workstation before or after the execution of a job, or to store data permanently in archives.

The previously described features already provide an effective tool to use resources of different computing centers both for capacity or capability computing, but many scientists and engineers use application packages. For applications without a graphical user interface, a tool kit simplifies the development of a custom built UNICORE plug-in. Over the years many plug-ins were developed, so that plug-ins already exist for many standard scientific applications, as e.g. for CPMD (Car-Parrinello Molecular Dynamics) [12], Fluent or MSC Nastran.

3.2. Server Tier

The server tier contains the Gateway and the Network Job Supervisor (NJS). The Gateway controls the access to a Usite and acts as the secure entry point accepting and authenticating UPL requests. A Usite identifies the participating organization (e.g. a supercomputing center) to the Grid with a symbolic name that resolves into the URL of the Gateway. An organization may be part of multiple Grids offering the same or different resources to different communities. The Gateway forwards incoming requests to the underlying Network Job Supervisor (NJS) of a virtual site (Vsite) for further processing. The NJS represents resources with a uniform user mapping scheme and no boundaries like firewalls between them.

A Vsite identifies a particular set of resources at a Usite and is controlled by a NJS. A Vsite may consist of a single supercomputer, e.g. a IBM p690 System with LoadLeveler, or a Linux cluster with PBS as resource management system. The flexibility of this concept supports different system architectures and gives the organization full control over its resources. Note that, there can be more than one Vsite inside each USite as depicted in Figure 2.

The NJS is responsible for the virtualization of the underlying resources by mapping the abstract job on a specific target system. This process is called “incarnation” and makes use of the Incarnation Database (IDB). System-specific data are stored in the IDB describing the software and hardware infrastructure of the system. Among others, the available resources like software, incarnation of abstract commands (standard UNIX command like rm, cp, ...) and site-specific administrative information are stored. In addition to the incarnation the NJS processes workflow descriptions included in an AJO, performs pre- and post-staging of files and authorizes the user via the UNICORE User Database (UUDB). Typically the Gateway and NJS are running on dedicated secure systems behind a firewall, although the Gateway could be placed outside a firewall or in a demilitarized zone.

3.3. Target System Tier

The Target System Interface (TSI) implements the interface to the underlying supercomputer with its resource management system. It is a stateless daemon running on the target system and interfacing with the local resource manager realized either by a batch system like PBS [13] or CCS [14], a batch system emulation on top of e.g. Linux, or a Grid resource manager like Globus’ GRAM [15,16].

3.4. Single Sign-On

The UNICORE security model relies on the usage of permanent X.509 certificates issued by a trusted Certification Authority (CA) and SSL based communication across ‘insecure’ networks. Certificates are used to provide a single sign-on in the client. The client unlocks the user’s keystore when it is first started, so that no further password requests are handed to the user. All authentication and authorization is done on the basis of the user certificate. At each UNICORE site user certificates are mapped to local accounts (standard UNIX uid/gid), which may be different at each site, due to existing naming conventions. The sites retain full control over the acceptance of users based on the identity of the individual – the distinguished name – or other information that might be contained in the certificate. UNICORE can handle multiple user certificates, i.e. it permits a client to be part of multiple, disjoint Grids. It is also possible to specify project accounts in the client allowing users to select different accounts for different projects on one execution system or to assume different roles with different privileges.

The private key in the certificate is used to sign each job and all included sub-jobs during the transit from the client to sites and between sites. This protects against tampering while the job is transmitted over insecure internet connections and it allows to verify the identity of the owner at the receiving end, without having to trust the intermediate sites which forwarded the job.

4. UNICORE Based Projects

During the evolutionary development of the UNICORE technology, many European and international projects have decided to base their Grid software implementations on UNICORE or to extend the growing set of core UNICORE functions with new features specific to their project focus. The goals and objectives of projects using UNICORE are not limited to the computer science community alone. Several other scientific domains such as bio-molecular engineering or computational chemistry are using the UNICORE technology as the basis of their work. In the following we present short overviews of goals and objectives of UNICORE-based projects and describe additional functions and services contributed to the UNICORE development.

4.1. EUROGRID – Application Testbed for European Grid Computing

In the EUROGRID⁴ project [17] a Grid network of leading European High Performance Supercomputing centers was established. Based on the UNICORE technology application-specific Grids were integrated, operated and demonstrated:

- Bio-Grid for biomolecular science
- Meteo-Grid for localized weather prediction
- CAE-Grid for coupling applications
- HPC-Grid for general HPC end-users

⁴funded by EC grant IST-1999-20247, duration: November 2000 - January 2004

As part of the project, the UNICORE software was extended by an efficient data transfer mechanism, resource brokerage mechanisms, tools and services for Application Service Providers (ASP), application coupling methods, and an interactive access feature [18]. Efficient data transfer is a important issue, as Grids typically rely on public parts of Internet connections. The available limited bandwidth has to be used efficiently to reduce the transfer time and the integrity of the transferred data has to be maintained, even if the transfer is interrupted. Depending on the application domain, additional security and confidentiality concerns need to be considered. This UNICORE high performance data transfer also uses X.509 certificates for authentication and encryption. To achieve not only a fast and secure transfer of data, but also high-performance capabilities, network Quality of Service (QoS) aspects, overlapping of streamed data transfers, and packet assembling and compression techniques are included.

In order to optimize the selection of resources – either done by the users manually or by a metascheduler automatically – resource brokerage mechanisms and detailed resource description abilities are important. Within the EUROGRID project, mechanisms were added to UNICORE, which allow users to specify their jobs in an abstract way improving the overall resource selection and accounting. In particular for the benefit of the industrial user aspects of security, convenience, and cost efficiency were addressed. To this end, the already existing security concepts of UNICORE were thoroughly evaluated and assessed as being adequate, hence no additional development had to be done. The task of the developed resource broker is to match the abstract specification of the users jobs and their requirements with the available resources in the Grid. The resource broker reports the best match back to the user including an estimate of the costs, which then allows the user to assign the appropriate resources to the job. For the suppliers of Grid resources (e.g. supercomputing centers) the resource broker allows to specify information about computational resources, architectures, processing power, storage and archiving facilities, post-processing facilities like visualization equipment, available software packages, and security guarantees. All this data is enhanced by billing information.

Supercomputing centers converge from pure providers of raw supercomputing power to Application Service Providers (ASP) running relevant scientific applications. For accounting and billing purposes the ASP needs to know the exact resources consumed by each customer in each run. For measuring the usage of supercomputers standard mechanisms provided by the resource management and operating system can be used, but measuring the usage of licenses requires a sophisticated approach. For some applications, e.g. from the Computer Aided Engineering (CAE) domain, this includes a connection to the applications licence manager. Establishing a link to the above mentioned resource broker is required to influence their decisions.

For solving complex problems applications from different domains, e.g. fluid-structure or electromagnetism-structure, need to be coupled. This is established by using the EUROGRID resource broker functionality and combining it with the available Metacomputing functionality developed in the UNICORE Plus project (cf. Section 1), which allows different schedulers of compute and application resources to cooperate. Finally, an interactive access to control and steer running application is needed for many scientific applications. The interactive use includes an interactive shell to actually login to computing resources using the UNICORE technology and security infrastructure.

EUROGRID used the UNICORE technology to provide the above described services and functionalities by developing new components. After the project ended, the developed components were revised and useful additions to the core UNICORE functions are now part of the available UNICORE software.

4.2. GRIP – Grid Interoperability Project

Grid computing empowers users and organizations to work effectively in an information-rich environment. Different communities and application domains have developed distinct Grid implementations some based on published open standards or on domain and community specific features. GRIP⁵ [19] had the objective to demonstrate that the different approaches of two distinct grids can successfully complement each other and that different implementations can interoperate. Two prominent Grid systems were selected for this purpose: UNICORE and GlobusTM[20], a toolkit developed in the United States. In contrast to UNICORE, Globus provides a set of APIs and services which requires more in-depth knowledge from the user. Globus is widely used in numerous international projects and many centers have Globus installed as Grid middleware.

The objectives of GRIP were:

- Develop software to enable the interoperation of independently developed Grid solutions
- Build and demonstrate prototype inter-Grid applications
- Contribute to and influence international Grid standards

During the runtime of the GRIP project the Open Grid Service Architecture was proposed by the Global Grid Forum (GGF) [21]. The arrival of OGSA also was an opportunity to influence the standards directly which were to be created and to start developments that allow UNICORE to interoperate not only with Globus but with services on the Grid in general, once the definition of the services and their interfaces became mature. OGSA did not change the overall objectives of GRIP, however, it influenced directly some of the technical results.

A basic requirement of GRIP was that the Grid interoperability layer should not change the well-known UNICORE user environment. As developers from both communities co-operated in the GRIP project, this goal was reached with only little changes of the UNICORE server components and no changes of the Globus Toolkit. This was achieved by the development of the so called Globus Target System Interface (Globus TSI), which provides UNICORE-access to computational resources managed by Globus. The Globus TSI was integrated into a heterogeneous UNICORE and Globus testbed.

To achieve the main objective of GRIP, the interoperability between UNICORE and Globus and initial OGSA services, the following elements had to be implemented:

- The interoperability layer between UNICORE and Globus Version 2
- The interoperability layer between UNICORE and Globus Version 3

⁵funded by EC grant IST-2001-32257, duration: January 2002 - February 2004

- The Access from UNICORE to simple Web services as a first step towards full integration of Web services
- The Interoperability of the certificate infrastructures of UNICORE and Globus
- A resource broker capable of brokering between UNICORE and Globus resources
- The Ontology of the resource description on an abstract level

In GRIP, two important application areas were selected to prove that the interoperability layers work as specified:

- Bio-molecular applications were instrumented in such a way that they are Grid-aware in any Grid environment and capable to seamlessly use UNICORE and Globus managed resources. The techniques developed in GRIP were designed and implemented in a generalized way to ensure that they can be used in other application domains as well.
- A meteorological application, the Relocatable Local Model (RLM), was decomposed in such a way that the components could execute on the most suitable resources in a Grid, independent of the middleware.

The results of the GRIP project are important for understanding general interoperability processes between Grid middleware systems. The experience and knowledge of the GRIP partners allowed to work in many relevant areas within GGF, like security, architecture, protocols, workflow, production management, and applications, and to influence the work in GGF.

4.3. OpenMolGRID – Open Computing Grid for Molecular Science and Engineering

The OpenMolGRID⁶ project [22] was focused on the development of Grid enabled molecular design and engineering applications. *In silico* testing [23] has become a crucial part in the molecular design process of new drugs, pesticides, biopolymers, and biomaterials. In a typical design process $O(10^5)$ to $O(10^6)$ candidate molecules are generated and their feasibility has to be tested. It is not economical to carry out experimental testing on all possible candidates. Therefore, computational screening methods provide a cheap and cost effective alternative to reduce the number of candidates. Over the years Quantitative Structure Activity/Property Relationship (QSAR/QSPR) methods have been shown to be reliable for the prediction of various physical, chemical, and biological activities [24].

QSPR/QSAR relies on the observation that molecular compounds with similar structure have similar properties. For each specific application a set of molecules is needed for which the target property is known. This requires searching globally distributed information resources for appropriate data. For the purpose of exploring molecular similarity, descriptors are calculated from the molecular structure. Thousands of molecular descriptors have been proposed and are used to characterize molecular structures with respect to different properties. Their calculation puts high demands on computer resources and requires high-performance computing.

⁶funded by EC grant IST-2001-37238, duration: September 2002 - February 2005

Based on this complex application the objectives of the OpenMolGRID project were defined as:

- Development of tools for secure and seamless access to distributed information and computational methods relevant to molecular engineering within the UNICORE frame
- Provision of a realistic testbed and reference application in life science
- Development of a toxicity prediction model validated with a large experimental set
- Provision of design principles for next generation molecular engineering systems.

In particular this included to use UNICORE to automatize, integrate, and speed-up the drug discovery pipeline.

The OpenMolGRID project addressed the objectives above by defining abstraction layers for data sources (databases) and methods (application software), and integrating all necessary data sources (e.g. ECOTOX [25]) and methods (e.g. 2D/3D Molecular Structure Conversion and Optimization, Descriptor Calculation, Structure Enumeration) into UNICORE. The project developed application specific user interfaces (plug-ins) and a mechanism to generate a complete UNICORE Job from an XML workflow specification. This so called Meta-Plug-in takes care of including all auxiliary steps like data format transformation and data transfers into the job, distributing data parallel tasks over available computational resources, and allocating resources to the tasks. Thereby the molecular design process was significantly improved as the time to build QSAR/QSPR models, the probability for mistakes, and the variability of results was reduced. In addition a command line client (CLC) for UNICORE was developed to enable the data warehouse to use Grid resources for its data transformation processes. The CLC offers the generation of UNICORE jobs from XML workflow description as well as the job submission, output retrieval, status query, and job abortion. The CLC consists of commands, an API, and a queuing component.

Besides the technical achievements of OpenMolGRID and the added value for pharmaceutical companies its results will contribute to the standardization of QSAR models.

4.4. VIOLA – Vertically Integrated Optical Testbed for Large Applications

The aim of the VIOLA⁷ project [26] is to build up a testbed with the latest optical network technology (multi 10 Gigabit Ethernet links). The goals and objectives of VIOLA are:

- Testing of new network components and network architectures
- Development and testing of software for dynamic bandwidth management
- Interworking of network technology from different manufacturers
- Development and testing of new applications from the Grid and Virtual Reality (VR) domain

⁷funded by BMBF grant 01AK605F, duration: May 2004 - April 2007

and the Meta-Scheduling plug-in generates the necessary scheduling information. The follow-on generation of the Meta-Scheduling framework will then be tightly integrated within UNICORE/GS (cf. Section 5.1.1).

4.5. NaReGI – National Research Grid Initiative

The Japanese NaReGI project [30] includes the UNICORE technology as the basic middleware for research and development. NaReGI is a collaboration project between industry, academia, and government. The goals and objectives are:

- Establishment of a national Japanese research Grid infrastructure
- Revitalization of the IT industry through commercialization of Grid middleware and strengthened international competitiveness
- Dissemination of Grid environments throughout industry
- Trailblazing the standardization of Grid technology
- Cultivation of human resources specializing in IT technology for Grids

Similar to the GRIP project (cf. Section 4) where an interoperability layer between UNICORE and Globus Toolkit 2 and 3 was developed, the NaReGI project plans to implement such a layer between UNICORE and Condor [31], called UNICONDOR. This interoperability layer will allow to submit jobs from the UNICORE client to Condor pools and to use Condor commands to submit jobs to UNICORE managed resources.

In the first phase of the NaReGI testbed UNICORE provides access to about 3000 CPUs in total with approximately 17 TFlops of peak performance. It is expected to increase the integrated peak performance to 100+ TFlops by the end of the project in 2007.

5. Future of UNICORE

The current UNICORE software implements a vertically integrated Grid architecture providing seamless access to various resources. Every resource is statically integrated into the UNICORE Grid by providing an interface to the appropriate resource manager.

One of the benefits Web services will bring to Grid computing is the concept of loosely coupled distributed services. Merging the idea of “everything being a service” with the achievements of the Grid community led to Grid services, enabling a new approach to the design of Grid architectures. The adoption of XML and the drive for standardization of the Open Grid Service Architecture provide the tools to move closer to the promise of interoperable Grids. A demonstrator validated the correspondence of UNICORE’s architectural model with the OGSA/OGSI (Open Grid Service Infrastructure [32]) approach, which encouraged the development of an OGSA/OGSI compliant UNICORE Grid architecture in the GRIP project (cf. Section 4.2).

In [16] UNICORE is examined for the evolution of a Grid system towards a service oriented Grid, primarily focussing on architectural concepts and models. Based on the current architecture and the enhancements provided by GRIP, first steps already integrate Web services into UNICORE. This included the provision of OGSI compliant port types

parallel to the proprietary ones as well as the design of XML based protocols. This work was continued in the UniGrids project.

As mentioned above the development of a Grid middleware is an continuous process of integrating new features, services, and adapting to emerging standards, and UNICORE is no exception. In the following we present new developments, some technical details, and report on projects, which enhance the UNICORE technology to serve the demands of the Grid in the future [33].

5.1. UniGrids – Uniform Interface to Grid Services

The strength of the UNICORE architecture is well-proven as described above. The rapid definition and adoption of OGSA allow the UNICORE development community to re-cast and extend the concepts of UNICORE through the use of Web services technologies. The goal of the UniGrids⁸ project [34] is to lift UNICORE on an architecture of loosely-coupled components while keeping its 'end-to-end' nature.

Thus, the integration of Web services techniques and UNICORE, which already started in the GRIP project (cf. Section 4.2), will continue in the UniGrids project. Interoperability, through adopting and influencing standards, form the philosophical foundation for UniGrids. The project aims to transform UNICORE into a system with interfaces that are compliant with the Web Services Resource Framework (WS-RF) [35] and that interoperate with other WS-RF compliant software components.

Such an approach offers great advantages both for the ease of development of new components by aggregation of services and through the integration of non-UNICORE components into the standards-based infrastructure.

In this sense, work is continuing in the following areas:

- Development of a compliant WS-RF hosting environment used for publishing UNICORE job and file services as Web services.
- Support of dynamic virtual organizations by enhancing the UNICORE security infrastructure to allow different usage models such as delegation and collective authorization.
- Development of translation mechanisms, such as resource ontologies, to interoperate with other OGSA compliant systems. Support for Grid economics by developing a Service Level Agreement (SLA) framework and cross-Grid brokering services.
- Development and integration of generic software components for visualization and steering of simulations (VISIT [36]), device monitoring and control, and tools for accessing distributed data and databases.

Applications from the scientific and industrial domain, like biomolecular and computational biology, geophysical depth imaging by oil companies, automotive, risk-management, energy, and aerospace are used to prove the developments in UniGrids.

The development in the UniGrids project will lead to UNICORE/GS, which follows the architecture of OGSA through the standardization of WS-RF and related work like e.g. the Web Services Notification technology [37]. The results will be made available under an open source BSD license.

⁸funded by EC grant IST-2002-004279, duration: July 2004 - June 2006

5.1.1. UNICORE/GS

Web service technology, and in particular the WS-RF, forms the basis for the UNICORE/GS software. WS-RF is the follow-on to OGSI, but more in line with mainstream Web services architecture [38]. Based on this new technology, UNICORE/GS will retain its key characteristics of seamlessness, security, and intuitiveness from both the user and administrative perspective, but will be built on a service oriented framework. This means that there is a loosening of the coupling between the components of the system. UNICORE/GS keeps the classical UNICORE topology of Usites, each containing a number of Vsites, but provides a new framework for integrating other services and providing common infrastructure functionality as services. This has the implication that new services will be easily integrated into the UNICORE/GS environment. Conversely, UNICORE/GS will be well-prepared to make use of external services.

The WS-RF technology is used to model core functionalities such as job submissions and file transfers as WS-Resources. These services are accessible via web service interfaces and thus establishing the UniGrids atomic services layer. This layer will be realized making extensive use of existing UNICORE server components.

All services in a Usite are accessible through the UniGrids Gateway that provides a secure entrance into the UNICORE/GS infrastructure. The principal is exactly the same as for classic UNICORE, however, the Gateway now routes messages according to Web Services Addressing (WS-Addressing) [39]. Authentication is based on transport level HTTPS security, although the intention is to move to Web Services Security (WS-Security) [40]. Regarding authorized access to resources, the UNICORE User Database (UUDB) will be available as a service to other services in the Usite, and will form the basis for future work concerning virtual organizations and fine-grained authorization schemes.

The underlying UniGrids atomic services layer will provide an excellent framework to deploy higher-level services such as co-allocation schedulers, workflow engines, and services for provision and easy access to data-intensive, remotely-steerable simulations.

5.2. NextGrid – Architecture for Next Generation Grids

In comparison to the UniGrids project which evolves the existing UNICORE Grid system to a service-oriented one, the NextGRID⁹ [41] project aims for the future: The goal is to provide the foundations for the next generation of Grids. NextGRID is not a project based on the UNICORE architecture or Grid system as-is, but institutions and people involved in the UNICORE development from the beginning on contribute expertise and experience to NextGRID.

Since it is obvious that there is no such thing as the one and only next generation Grid, and experts envisage the co-existence of multiple Grids with well-defined boundaries and access points, NextGRID is going to define a Grid architecture which can be seen as building blocks for Grids. It does not only provide interoperability by-design between entities which exist within one instantiation of such an architecture, but it also facilitates the interoperability between different Grids developed according to the NextGRID architecture.

Although developing a Grid one generation ahead, NextGRID is not starting from scratch. Properties to incarnate and functions to realize future Grids are expertly de-

⁹funded by EC grant IST-2002-511563, duration: September 2004 - August 2007

scribed in [42] and [33]. These reports frame NextGRID's architectural development while the Open Grid Services Architecture is going to define Grid services and their interactions and does therefore make up a starting point for the conceptualization and design of NextGRID. In addition, regarding the underlying technology and architectural model, NextGRID propagates the usage of Web Services and the adoption of Service-Oriented Architecture (SOA) [43] concepts and models.

NextGRID focuses on security, economic sustainability, privacy/legacy, scalability and usability. The following properties have the highest priorities when carrying out the following work:

- Developing an architecture for next generation Grids
- Implementing and testing prototypes aligned with the concepts and design of the NextGRID architecture
- Creating reference applications which make use of the NextGRID prototypes
- Facilitating the transition from scientific- to business-oriented Grids by integrating the means to negotiate a certain Quality of Service (QoS) level
- Specifying the methods, processes, and services necessary to dynamically operate Grids across multiple organizations which comprise heterogeneous resources

Since the ongoing UNICORE development in projects like UniGrids shares resources as well as the technological foundation with NextGRID there is a high chance that the outcome of NextGRID will also represent the next step of UNICORE's evolution.

6. Conclusion

In this paper we presented the evolution of the UNICORE technology from a Grid software with prototype character developed in two German projects to a full-grown, well-tested, widely used and accepted Grid middleware. UNICORE – Uniform Interface to Computing Resources – provides a *seamless, secure and intuitive* access to distributed Grid resources. Although the UNICORE vision was already coined in 1997, the then stated goals and objectives of hiding the seams of resource usage, incorporating a strong security model, and providing an easy to use graphical user interface for scientists and engineers are still valid today: to achieve these goals and objectives, UNICORE is designed as a vertically integrated Grid middleware providing components at all layers of a Grid infrastructure, from a graphical user interface down to the interfaces to target machines.

Initially developed in the German projects UNICORE and UNICORE Plus, UNICORE was soon established as a promising Grid middleware in several European projects. In the GRIP project an interoperability layer between UNICORE and the Globus Toolkit 2 and 3 was developed to demonstrate the interoperability of independently developed Grid solutions, allowing to build and to demonstrate inter-Grid applications from the bio-molecular and meteorological domain. In the EUROGRID project, European high performance supercomputing centers joined to extend UNICORE with an efficient data transfer, resource brokerage mechanisms, ASP services, application coupling methods,

and an interactive access. In addition, a Bio-Grid, Meteo-Grid, CAE-Grid, and HPC-Grid were established to integrate a variety of application domains. The main objective of the OpenMolGRID project is to provide a unified and extensible information-rich environment based on UNICORE for solving problems from molecular science and engineering. In the VIOLA project a vertically integrated testbed with the latest optical network technology is built up. UNICORE is used as the Grid middleware for enabling the development and testing of new applications in the optical networked testbed, which provides advanced bandwidth management and QoS features.

With these developments UNICORE grew to a software system usable in production Grids. In this context UNICORE is deployed in the large German supercomputing centers to provide access to their resources. At the John von Neumann-Institute for Computing, Research Center Jülich, many users submit their batch jobs through UNICORE to the 1312-processor 8.9 TFlop/s IBM p690 cluster and the Cray SV1 vector machine. Leading European HPC centers joined in the project DEISA to build a distributed European infrastructure for scientific applications based on UNICORE to build and operate a distributed multi tera-scale supercomputing facility.

The future of UNICORE is promising and follows the trend of “Everything being a Service” by adapting to Open Grid Service Architecture (OGSA) standards. In this context, the UniGrids project continues the effort of the GRIP project in integrating the Web Services and UNICORE technology to enhance UNICORE to an architecture of loosely-coupled components while keeping its “end-to-end” nature. To this end UNICORE/GS will be developed, which makes UNICORE compliant with the Web Services Resource Framework (WS-RF).

Today the UNICORE software is available as open source under a BSD licence from SourceForge for download. This enables the community of core UNICORE developers to grow and makes future development efforts open to the public.

7. Acknowledgments

The work summarized in this paper was done by many people. We gratefully thank them for their past, present, and future contributions in developing the UNICORE technology. Most of the work described here was supported and funded by BMBF and different programmes of the European Commission under the respective contract numbers mentioned above.

REFERENCES

1. I. Foster, C. Kesselman (Eds.). *The Grid: Blueprint for a New Computing Infrastructure*. Morgan Kaufmann Publishers Inc. San Fransisco, 1999.
2. I. Foster and C. Kesselman. Globus: A Metacomputing Infrastructure Toolkit. *International Journal on Supercomputer Applications*, 11(2):115–128, 1997.
3. D. Erwin (Ed.). *UNICORE - Uniformes Interface für Computing Ressourcen, final project report (in German)*. 2000.
4. D. Erwin (Ed.). *UNICORE Plus Final Report - Uniform Interface to Computing Resources*. Forschungszentrum Jülich, 2003.
5. UNICORE at SourceForge. <http://unicore.sourceforge.net>.

6. Jump - Juelich Multi Processor, 8,9 TFlop/s IBM p690 eServer Cluster. <http://jumpdoc.fz-juelich.de>.
7. DEISA - Distributed European Infrastructure for Supercomputing Applications. <http://www.deisa.org>.
8. M. Romberg. The UNICORE Grid Infrastructure. *Scientific Programming*, 10(2):149–157, 2002.
9. I. Foster, C. Kesselmann, J. M. Nick, and S. Tuecke. The Physiology of the Grid. In F. Berman, G. C. Fox, and A. J. G. Hey, editors, *Grid Computing*, pages 217–249. John Wiley & Sons Ltd, 2003.
10. D. Snelling, S. van den Berghe, G. von Laszweski, Ph. Wieder, D. Breuer, J. MacLaren, D. Nicole, and H.-Ch. Hoppe. A UNICORE Globus Interoperability Layer. *Computing and Informatics*, 21:399–411, 2002.
11. D. Snelling. UNICORE and the Open Grid Services Architecture. In F. Berman, G. Fox, and T. Hey, editor, *Grid Computing: Making The Global Infrastructure a Reality*, pages 701–712. John Wiley & Sons, 2003.
12. V. Huber. Supporting Car-Parrinello Molecular Dynamics Application with UNICORE. In *Proc. of the International Conference on Computational Science (ICCS 2001)*, pages 560–566, 2001.
13. PBS - Portable Batch System. <http://www.openpbs.org/>.
14. M. Hovestadt, O. Kao, A. Keller, and A. Streit. Scheduling in HPC Resource Management Systems: Queuing vs. Planning. In D. G. Feitelson and L. Rudolph, editor, *Proc. of the 9th Workshop on Job Scheduling Strategies for Parallel Processing*, volume 2862 of *Lecture Notes in Computer Science*, pages 1–20. Springer, 2003.
15. Globus: Research in Resource Management. <http://www.globus.org/research/resource-management.html>.
16. R. Menday and Ph. Wieder. GRIP: The Evolution of UNICORE towards a Service-Oriented Grid. In *Proc. of the 3rd Cracow Grid Workshop (CGW'03)*, pages 142–150, 2003.
17. EUROGRID - Application Testbed for European Grid Computing. <http://www.eurogrid.org>.
18. D. Mallmann. EUROGRID - Application Testbed for European Grid Computing. In *Proc. of Industrial GRID Conference 2001*, 2001.
19. GRIP - Grid Interoperability Project. <http://www.grid-interoperability.org>.
20. The Globus Project. <http://www.globus.org/>.
21. Global Grid Forum. www.globalgridforum.org.
22. OpenMolGRID - Open computing Grid for Molecular Science and Engineering. <http://www.openmolgrid.org>.
23. S. Sild, U. Maran, M. Romberg, B. Schuller, and E. Benfenati. OpenMolGRID: Using Automated Workflows in GRID Computing Environment. In *Proceedings of the European Grid Conference 2005*, 2005.
24. M. Karelson. *Molecular Descriptors in QSAR/QSPR*. John Wiley & Sons, New York, 2000.
25. U.S. Environmental Protection Agency ECOTOXicology database. <http://www.epa.gov/ecotox/>.
26. VIOLA - Vertically Integrated Optical Testbed for Large Applications in DFN. [http:](http://)

- [//www.viola-testbed.de](http://www.viola-testbed.de).
27. MetaMPICH - Flexible Coupling of Heterogeneous MPI Systems. <http://www.lfbs.rwth-aachen.de/~martin/MetaMPICH/>.
 28. G. Quecke and W. Ziegler. MeSch - An Approach to Resource Management in a Distributed Environment. In R. Buyya, editor, *Proc. of 1st IEEE/ACM International Workshop on Grid Computing (Grid 2000)*, volume 1971 of *Lecture Notes in Computer Science*, pages 47–54. Springer, 2000.
 29. Global Grid Forum GRAAP (Grid Resource Allocation Agreement Protocol) Working Group. <https://forge.gridforum.org/projects/graap-wg/>.
 30. NaReGI - National Research Grid Initiative. http://www.naregi.org/index_e.html.
 31. The Condor Project. <http://www.cs.wisc.edu/condor/>.
 32. S. Tuecke, K. Czajkowski, I. Foster, J. Frey, S. Graham, C. Kesselman, T. Maquire, T. Sandholm, D. Snelling, and P. Vanderbilt (eds.). The Open Grid Services Infrastructure (OGSI) Version 1.0, 2003.
 33. Next Generation Grids 2 – Requirements and Options for European Grids Research 2005-2010 and Beyond. ftp://ftp.cordis.lu/pub/ist/docs/ngg2_eg_final.pdf, 2004.
 34. UniGrids - Uniform Access to Grid Services. <http://www.unigrids.org>.
 35. OASIS Web Services Resource Framework (WSRF). <http://www.oasis-open.org/committees/wsrf>.
 36. VISIT – A Visualization Toolkit. <http://www.fz-juelich.de/zam/visit/>.
 37. OASIS Web Services Notification (WSN). <http://www.oasis-open.org/committees/wsn>.
 38. W3C Web Services Architecture. <http://www.w3.org/TR/ws-arch/>.
 39. Web Services Addressing. <http://www.w3.org/Submission/2004/SUBM-ws-addressing-20040810/>.
 40. OASIS Web Services Security (WSS). <http://www.oasis-open.org/committees/wss/>.
 41. NextGRID - Architecture for Next Generation Grids. <http://www.nextgrid.org>.
 42. Next Generation Grids – European Grid Research 2005 - 2010. ftp://ftp.cordis.lu/pub/ist/docs/ngg_eg_final.pdf, 2003.
 43. T. Erl. *Service-Oriented Architecture: A Field Guide to Integrating XML and Web Services*. Prentice Hall PTR, 2004.